

# Packages

To run the following code, make sure you have loaded (and installed) the following packages. We like to set **ggplot2** to use the minimal theme, but this is of course entirely optional.

```
library(cleanNLP)
library(dplyr)
library(readr)
library(stringi)
library(ggplot2)
library(tokenizers)
```

```
## Warning: package 'tokenizers' was built under R version 3.4.4
```

```
theme_set(theme_minimal())
```

## Splitting text into words

Consider a string in R containing the first paragraph of text from the novel *L'Étranger* of Albert Camus (we'll use `stri_wrap` just to fit the output on the slide):

```
stri_wrap(letranger)
```

```
## [1] "Aujourd'hui, maman est morte. Ou peut-être hier, je ne"  
## [2] "sais pas. J'ai reçu un télégramme de l'asile: «Mère décédée."  
## [3] "Enterrement demain. Sentiments distingués.» Cela ne veut rien"  
## [4] "dire. C'était peut-être hier."
```

In order to work with this text, a good first step is to split it apart into its constituent words.

## Splitting with whitespace

Splitting on whitespace alone works reasonably well, though there are some issues with punctuation marks:

```
stri_split(letranger, fixed = " ")
```

```
## [[1]]
##  [1] "Aujourd'hui,"      "maman"           "est"
##  [4] "morte."            "Ou"              "peut-être"
##  [7] "hier,"             "je"              "ne"
## [10] "sais"              "pas.J'ai"        "reçu"
## [13] "un"                "télégramme"      "de"
## [16] "l'asile:"          "«Mère"           "décédée."
## [19] "Enterrement"       "demain.Sentiments" "distingués.»"
## [22] "Cela"              "ne"              "veut"
## [25] "rien"              "dire."           "C'était"
## [28] "peut-êtrehier."
```

## Splitting with cleanNLP

There are a number of packages that support the more complex logic needed to deal with many of these errors. Here we'll use the **cleanNLP** package as it will be easy to adapt our approach to work with more complex annotators in the next section.

## Running annotations

We start by initialising the tokenizers back end within the **cleanNLP** package. We'll indicate that we want a French locale as this input text is in French.

```
library(cleanNLP)
init_tokenizers(locale = "fr")
```

Then, we run the annotators over the text. We set the option `as_strings` because we are passing the text into the function as a raw string:

```
letranger_anno <- run_annotators(letranger, as_strings = TRUE)
letranger_anno
```

```
##
## A CleanNLP Annotation:
##   num. documents: 1
```

## An annotation object

The result seems to be wrapped up in a fairly complex object; however, it is nothing more than a list of data frames. To collapse all of these lists into a one table summary of the tokenisation process, we will call the function `get_combine` on the annotation object:

```
letranger_tokens <- get_combine(letranger_anno)
```

## NOTE: `get_combine` has been renamed `cnlp_get_tif`

## An annotation object

The result is a data frame with one row for each token. Meta data about each token, such as the sentence number and character offset, are included as columns.

```
letranger_tokens
```

```
## # A tibble: 42 x 6
##   id      sid  tid word      cid spaces
##   <chr> <int> <int> <chr>    <int>  <dbl>
## 1 doc1     1     1 Aujourd'hui     1     0
## 2 doc1     1     2 ,             12    1.00
## 3 doc1     1     3 maman          14    1.00
## 4 doc1     1     4 est           20    1.00
## 5 doc1     1     5 morte          24     0
## 6 doc1     1     6 .             29    1.00
## # ... with 36 more rows
```

## cleanNLP tokenization results

Notice that the resulting tokens fix most of the problems in the original white space based technique:

```
letranger_tokens$word
```

##	[1]	"Aujourd'hui"	","	"maman"
##	[4]	"est"	"morte"	". "
##	[7]	"Ou"	"peut"	"_ "
##	[10]	"être"	"hier"	","
##	[13]	"je"	"ne"	"sais"
##	[16]	"pas.J'ai"	"reçu"	"un"
##	[19]	"télégramme"	"de"	"l'asile"
##	[22]	":"	"«"	"Mère"
##	[25]	"décédée"	". "	"Enterrement"
##	[28]	"demain.Sentiments"	"distingués"	". "
##	[31]	"»"	"Cela"	"ne"
##	[34]	"veut"	"rien"	"dire"
##	[37]	". "	"C'était"	"peut"
##	[40]	"_ "	"êtrehier"	". "



## Corpus Metadata

A dataset where each record is its own text is known as a *corpus*. In the remainder of this session, we will be working with a corpus of the 56 short stories featuring Sherlock Holmes.

We start by constructing a meta data table of these stories:

```
paths <- dir("data/holmes_stories", full.names = TRUE)
sh_meta <- data_frame(id = as.integer(seq_along(paths)),
                      story = stri_sub(basename(paths), 4, -5))
sh_meta %>% print(n = 5)
```

```
## # A tibble: 56 x 2
##       id story
##   <int> <chr>
## 1     1 a_scandal_in_bohemia
## 2     2 the_redheaded_league
## 3     3 a_case_of_identity
## 4     4 the_boscombe_valley_mystery
## 5     5 the_five_orange_pips
## # ... with 51 more rows
```

## Annotating files on disk

We want to construct a similar data frame of tokens for all of the short stories in our corpus. As a first step, we will re-initialise the tokenizers backend using an English locale:

```
library(cleanNLP)
init_tokenizers(locale = "en_GB")
```

Then, we call the annotation engine with the paths to the files instead of the raw text:

```
sh_anno <- run_annotators(paths)
```

And once again, collapse the object into a single table.

```
sh_tokens <- get_combine(sh_anno)
```

## NOTE: get\_combine has been renamed cnlp\_get\_tif

## Sherlock Holmes tokens

The resulting table, as before, has one row for each token in the original dataset.

```
library(magrittr)
sh_tokens %<>% mutate(id = as.integer(gsub("doc","",id)))
sh_tokens
```

```
## # A tibble: 550,697 x 6
##       id    sid    tid word      cid spaces
##   <int> <int> <int> <chr>   <int> <dbl>
## 1     1     1     1  To       2    1.00
## 2     1     1     2 Sherlock  5    1.00
## 3     1     1     3 Holmes   14    1.00
## 4     1     1     4 she     21    1.00
## 5     1     1     5 is      25    1.00
## 6     1     1     6 always  28    1.00
## # ... with 5.507e+05 more rows
```

## Sherlock Holmes tokens, sentence 10

```
sh_tokens %>% filter(id == 1) %>% filter(sid == 10) %>% print(n = 12)
```

```
## # A tibble: 35 x 6
##       id    sid  tid word      cid spaces
##   <int> <int> <int> <chr>    <int>   <dbl>
## 1     1     10     1 Grit     875    1.00
## 2     1     10     2 in      880    1.00
## 3     1     10     3 a       883    1.00
## 4     1     10     4 sensitive 885    1.00
## 5     1     10     5 instrument 895    0
## 6     1     10     6 ,       905    1.00
## 7     1     10     7 or      907    1.00
## 8     1     10     8 a       910    1.00
## 9     1     10     9 crack   912    1.00
## 10    1     10    10 in     918    1.00
## 11    1     10    11 one    921    1.00
## 12    1     10    12 of     925    1.00
## # ... with 23 more rows
```

# Using tokens

One can often learn a lot about a corpus by simply finding the occurrences of certain tokens or patterns of tokens within it.

For example:

- ▶ length of sentences
- ▶ number of citations
- ▶ presence of known characters
- ▶ count of hashtags in a tweet
- ▶ ratio of quotes/dialogue to raw text

# Visualising Watson and Holmes

Where do Watson and Holmes occur within each text?

```
sh_tokens %>%  
  group_by(id) %>%  
  mutate(percentage_loc = sid / max(sid)) %>%  
  filter(word %in% c("Watson", "Holmes")) %>%  
  ggplot(aes(percentage_loc, id)) +  
    geom_point(aes(color = word), alpha = 0.5) +  
    geom_hline(yintercept = 44)
```

# Visualising Watson and Holmes

```
## Warning: package 'viridis' was built under R version 3.4.4
```




## Average sentence length

By counting the number of periods, question marks, and exclamation marks, we can approximate the average length of each sentence in the text.

```
sh_tokens %>%  
  mutate(sentence_end = word %in% c(".", "?", "!")) %>%  
  group_by(id) %>%  
  summarize(mean_sentence_len = n() / sum(sentence_end)) %>%  
  ggplot(aes(id, mean_sentence_len)) +  
    geom_line() +  
    geom_point()
```



# Average sentence length



## Frequency tokens

In our first set of analyses, we used our prior knowledge to determine which tokens would be interesting to identify and tabulate.

Alternatively, we can let the corpus itself tell us which terms would be the most interesting. Often just knowing which terms are the most frequent, or occur in certain patterns, is interesting itself.

## Naïve approach

To begin, we can see what the most common words are across the corpus using the count function:

```
sh_tokens %>%  
  count(id, word, sort = TRUE)
```

```
## # A tibble: 105,460 x 3  
##       id word      n  
##   <int> <chr> <int>  
## 1     22 ,      827  
## 2     28 .      822  
## 3     37 ,      794  
## 4     40 .      778  
## 5     22 .      770  
## 6     28 ,      765  
## # ... with 1.055e+05 more rows
```

## Stop words

What is interesting about these top terms? Cynically, we might say very little. The problem is that the most common terms are simply punctuation marks.

If we look farther down the list, common function words such as 'the' and 'my' dominate the counts.

```
## # A tibble: 6 x 3
##       id word      n
##   <int> <chr> <int>
## 1     1 I      259
## 2     1 to     242
## 3     1 of     235
## 4     1 and    227
## 5     1 a      212
## 6     1 in     152
```

A popular way of dealing with this problem is to define a list of *stop words*, those tokens that are common enough to be thematically uninteresting.

## Stop words

We have included a simple list of stop words in the dataset for today, which you should read in with the following.

```
stopwords <- readLines("data/stopwords_en.txt")
sample(stopwords, 25L)
```

```
## [1] "opened"      "said"        "puts"        " "
## [5] "alone"       "ended"       "nor"         "by"
## [9] "use"         "rooms"       "cannot"      "can"
## [13] "together"    "it"          "almost"     "asks"
## [17] "their"       "greater"     "gives"       "against"
## [21] "presents"   "yourselves" "likely"      "shouldn't"
## [25] "beings"
```

Take a moment to look at some of the words. What parts of speech dominate the list?

## Most common non-stopwords

We will make use of the `top_n` function to select just the top 10 occurrences within each text. We also add a call to `left_join` to explicitly add the names of each story to the output:

```
sh_toptokens <- sh_tokens %>%  
  filter(!(tolower(word) %in% stopwords)) %>%  
  count(id, word, sort = TRUE) %>%  
  group_by(id) %>%  
  top_n(n = 10, n) %>%  
  left_join(sh_meta, by = "id")
```

## Tokens from 'A Scandal in Bohemia'

```
sh_toptokens %>% filter(id == 1) %>% print(n = Inf)
```

```
## # A tibble: 11 x 4
##       id word          n story
##   <int> <chr>      <int> <chr>
## 1     1  Holmes        48 a_scandal_in_bohemia
## 2     2 photograph    21 a_scandal_in_bohemia
## 3     3   King        17 a_scandal_in_bohemia
## 4     4  Majesty       16 a_scandal_in_bohemia
## 5     5   house       14 a_scandal_in_bohemia
## 6     6  little       14 a_scandal_in_bohemia
## 7     7   Adler       13 a_scandal_in_bohemia
## 8     8   door       13 a_scandal_in_bohemia
## 9     9   hand       13 a_scandal_in_bohemia
## 10    10  Irene       13 a_scandal_in_bohemia
## 11    11 minutes      13 a_scandal_in_bohemia
```

## Tokens from 'His Last Bow'

```
sh_toptokens %>% filter(id == 44) %>% print(n = Inf)
```

```
## # A tibble: 12 x 4
##       id word          n story
##   <int> <chr>      <int> <chr>
## 1     44 Von         38 his_last_bow
## 2     44 Bork        34 his_last_bow
## 3     44 Holmes       21 his_last_bow
## 4     44 Watson       17 his_last_bow
## 5     44 American     12 his_last_bow
## 6     44 German       12 his_last_bow
## 7     44 little       12 his_last_bow
## 8     44 country      11 his_last_bow
## 9     44 car          10 his_last_bow
## 10    44 papers       10 his_last_bow
## 11    44 safe         10 his_last_bow
## 12    44 secretary    10 his_last_bow
```



## Why characters?

One particular category that floats to the top of our lists of most frequent tokens are the main characters for each story. Identifying the people mentioned in a corpus of text has many applications, including:

- ▶ on social media it indicates trending issues
- ▶ in news articles, the people mentioned give a good clue as to what topics are being discussed (politics, food, culture, local events, ..)
- ▶ in fiction, as we have seen, the presence and absence of characters is a major indicator of plot arcs

## Proper nouns

Some of the most frequent non stop words in the texts refer to the names of the characters. How might we extract these directly?

```
sh_propn <- sh_tokens %>%  
  filter(!(tolower(word) %in% stopwords)) %>%  
  filter((tolower(word) != word)) %>%  
  count(id, word, sort = TRUE) %>%  
  group_by(id) %>%  
  top_n(n = 10, n) %>%  
  left_join(sh_meta, by = "id")
```

## Proper nouns from 'A Scandal in Bohemia'

```
sh_propn %>% filter(id == 1) %>% print(n = Inf)
```

```
## # A tibble: 11 x 4
##       id word          n story
##   <int> <chr>      <int> <chr>
## 1     1  Holmes      48 a_scandal_in_bohemia
## 2     2   King      17 a_scandal_in_bohemia
## 3     3 Majesty     16 a_scandal_in_bohemia
## 4     4  Adler      13 a_scandal_in_bohemia
## 5     5  Irene      13 a_scandal_in_bohemia
## 6     6 Briony      11 a_scandal_in_bohemia
## 7     7  Lodge      11 a_scandal_in_bohemia
## 8     8 Sherlock     11 a_scandal_in_bohemia
## 9     9  Bohemia       7 a_scandal_in_bohemia
## 10    10  Norton       7 a_scandal_in_bohemia
## 11    11  Street       7 a_scandal_in_bohemia
```

## Proper nouns from 'His Last Bow'

```
sh_propn %>% filter(id == 44) %>% print(n = Inf)
```

```
## # A tibble: 10 x 4
##       id word          n story
##   <int> <chr>      <int> <chr>
## 1    44 Von         38 his_last_bow
## 2    44 Bork        34 his_last_bow
## 3    44 Holmes       21 his_last_bow
## 4    44 Watson       17 his_last_bow
## 5    44 American     12 his_last_bow
## 6    44 German       12 his_last_bow
## 7    44 Altamont      8 his_last_bow
## 8    44 England       8 his_last_bow
## 9    44 Martha        7 his_last_bow
## 10   44 Baron         6 his_last_bow
```

## One top character

Of course, two major characters are always going to be Sherlock Holmes and John Watson. Let us remove them from the list, as well as any names shorter than four characters (these are usually honorific rather than names). We will then take the most mentioned character from each text:

```
holmes_watson <- c("Sherlock", "Holmes", "John", "Watson")
sh_topchar <- sh_tokens %>%
  filter(stri_length(word) > 4) %>%
  filter(!(word %in% holmes_watson)) %>%
  filter(!(tolower(word) %in% stopwords)) %>%
  filter((tolower(word) != word)) %>%
  count(id, word) %>%
  left_join(sh_meta, by = "id") %>%
  group_by(id) %>%
  top_n(n = 1, n)
```

## One top character, cont.

```
sh_topchar %>% filter(id %in% c(1, 45)) %>%  
  print(n = Inf)
```

```
## # A tibble: 2 x 4  
##       id word          n story  
##   <int> <chr>    <int> <chr>  
## 1     1 Majesty     16 a_scandal_in_bohemia  
## 2    45 Baron      17 the_illustrious_client
```

## Top name

Sometimes this works well, sometimes it picks up the right idea but not enough to really know who the character is (i.e., “Colonel” or “Majesty”), and sometimes it works very well. We will see in the next session how to do a better job of this using more advanced annotation engines.

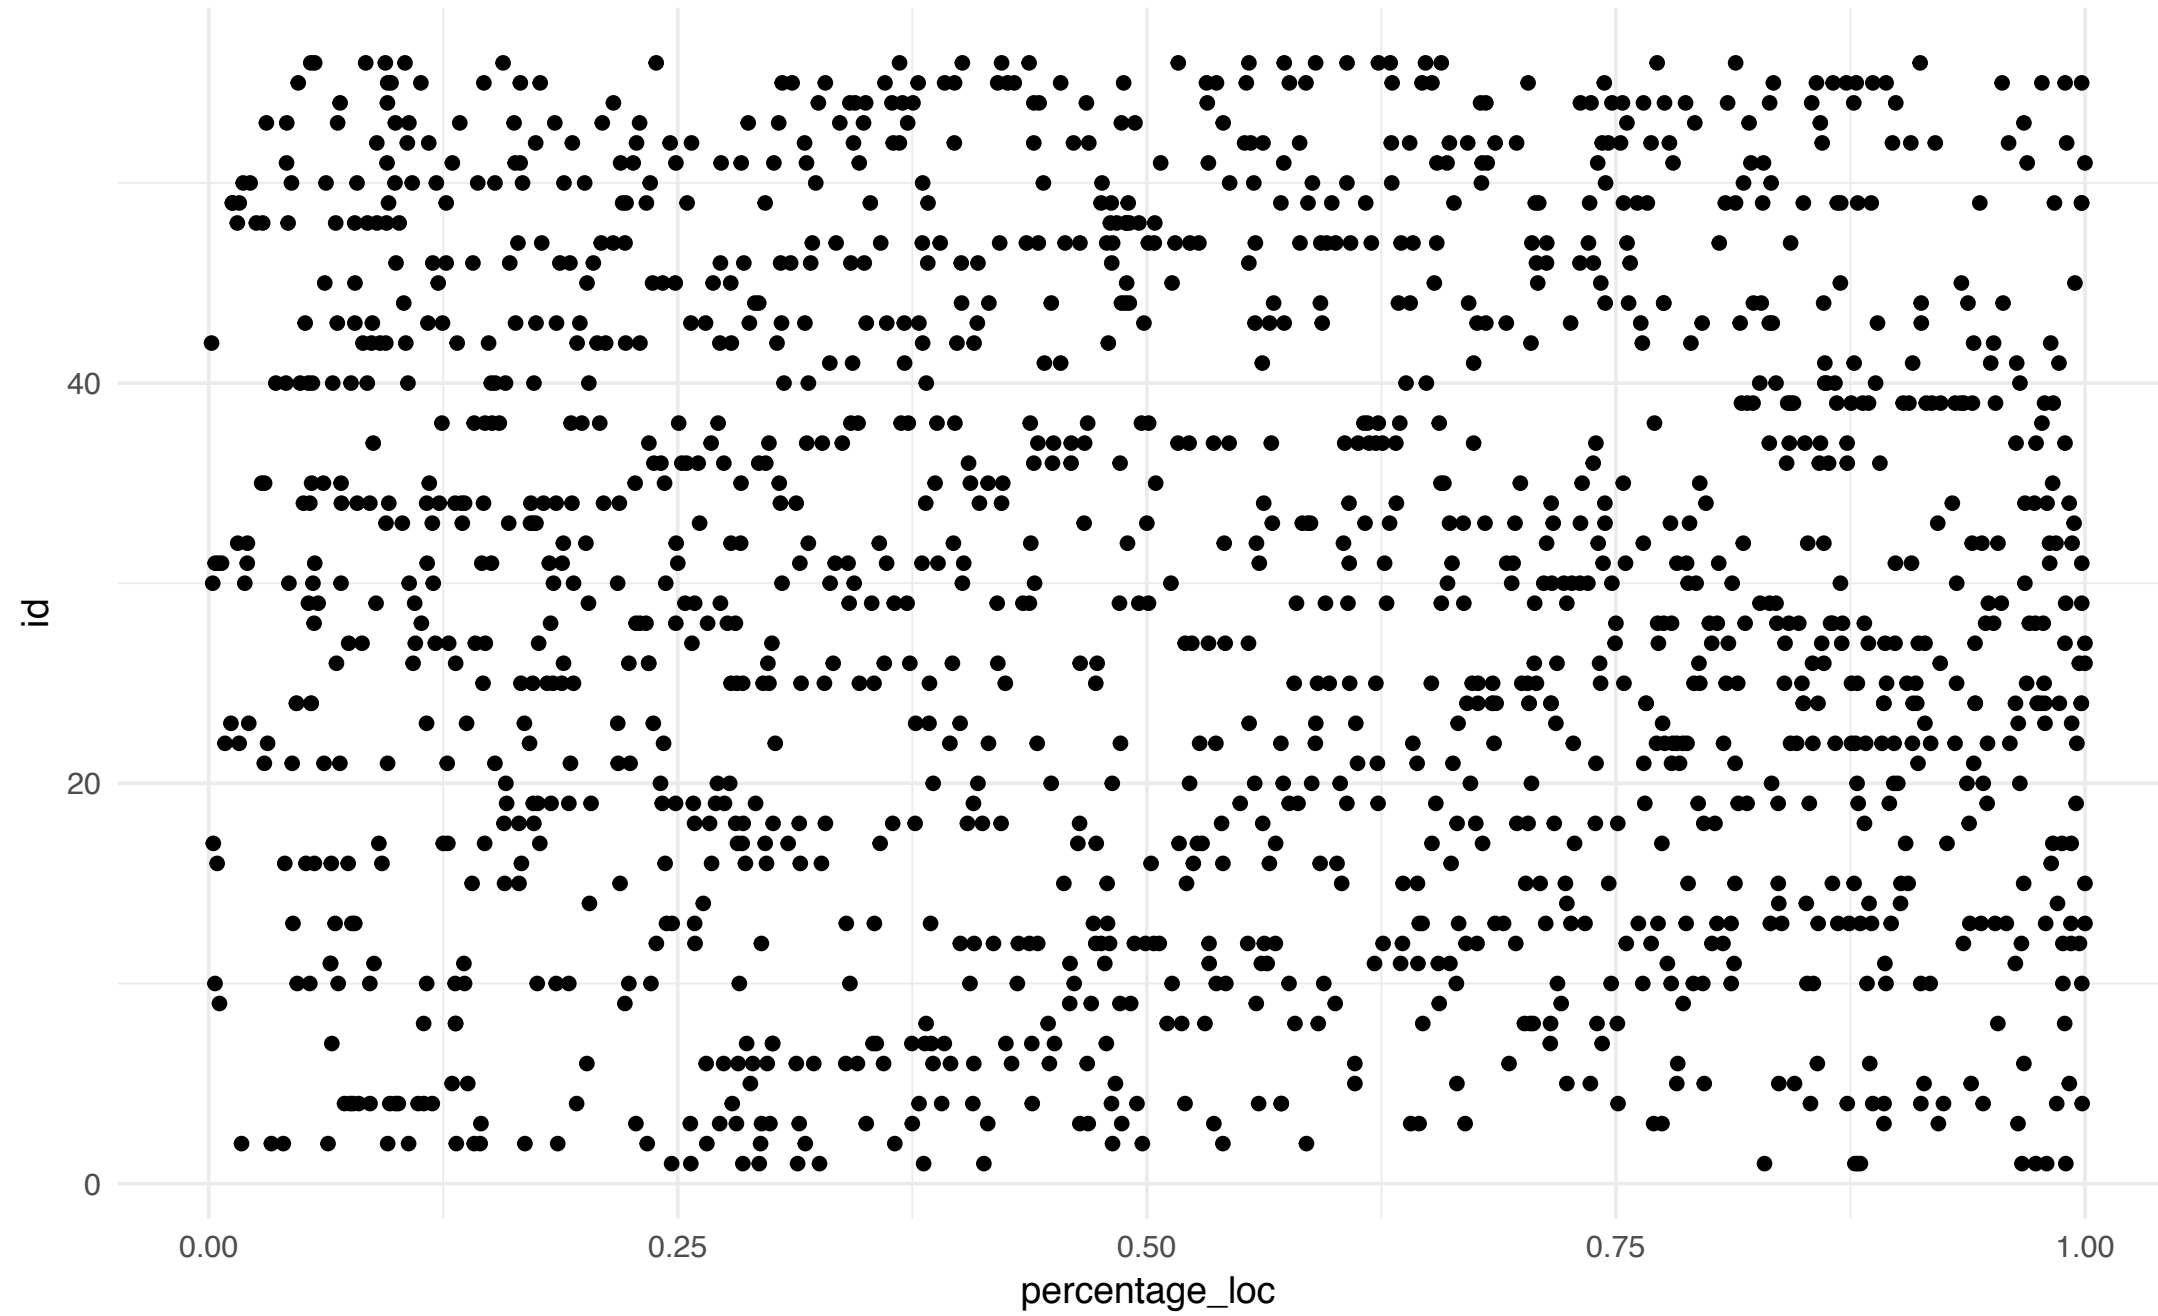
## Visualising main characters

We will make use of the `semi_join` function to plot the character locations:

```
sh_tokens %>%  
  group_by(id) %>%  
  mutate(percentage_loc = sid / max(sid)) %>%  
  semi_join(sh_topchar, by = c("id", "word")) %>%  
  ggplot(aes(percentage_loc, id)) +  
    geom_point()
```



## Visualising main characters



## Textual topics and themes

When we look back at our original list of top tokens, many of those instances that are not characters describe the main topics, themes, or artefacts of interest in the story.

Finding these frequent, non-proper nouns can indicate the theme or topics of interest within a corpus of texts.

## Non-proper words

Our original code can be easily modified to only count those with all lower-case letters:

```
sh_theme <- sh_tokens %>%  
  filter(!(tolower(word) %in% stopwords)) %>%  
  filter((tolower(word) == word)) %>%  
  count(id, word) %>%  
  group_by(id) %>%  
  top_n(n = 10, n) %>%  
  left_join(sh_meta, by = "id")
```

## Non-proper words, cont.

```
sh_theme %>% filter(id == 1) %>% print(n = 10)
```

```
## # A tibble: 10 x 4
```

```
##       id word          n story
```

```
##   <int> <chr>      <int> <chr>
```

```
## 1     1 door       13 a_scandal_in_bohemia
```

```
## 2     1 half       11 a_scandal_in_bohemia
```

```
## 3     1 hand       13 a_scandal_in_bohemia
```

```
## 4     1 house      14 a_scandal_in_bohemia
```

```
## 5     1 little     14 a_scandal_in_bohemia
```

```
## 6     1 matter     11 a_scandal_in_bohemia
```

```
## 7     1 minutes    13 a_scandal_in_bohemia
```

```
## 8     1 photograph 21 a_scandal_in_bohemia
```

```
## 9     1 street     11 a_scandal_in_bohemia
```

```
## 10    1 woman      12 a_scandal_in_bohemia
```

## Word frequencies

What we need is something stronger than a stop word list; conveniently such a dataset is included in the **cleanNLP** package as the dataset `word_frequency`.

```
word_frequency %>% print(n = 9)
```

```
## # A tibble: 150,000 x 3
##   language word   frequency
##   <chr>      <chr>     <dbl>
## 1 en        the       3.93
## 2 en        of        2.24
## 3 en        and        2.21
## 4 en        to        2.06
## 5 en        a         1.54
## 6 en        in         1.44
## 7 en        for        1.01
## 8 en        is         0.800
## 9 en        on         0.638
## # ... with 1.5e+05 more rows
```

## Filtering by word frequency

Instead of a stopwords list, we filter out those words with a certain frequency cut-off. By changing this tuning parameter, we can tweak the results until they look reasonable.

```
sh_wordfreq <- sh_tokens %>%  
  mutate(lemma = tolower(word)) %>%  
  inner_join(word_frequency, by = c("lemma" = "word")) %>%  
  filter(frequency < 0.01) %>%  
  filter((tolower(word) == word)) %>%  
  count(id, word) %>%  
  group_by(id) %>%  
  top_n(n = 10, n) %>%  
  left_join(sh_meta, by = "id") %>%  
  arrange(id, desc(n))
```

A more complex method could compare these global probabilities to the frequency in our text and identify the most deviant probabilities.

## Filtering by word frequency, cont.

```
sh_wordfreq %>% filter(id == 1) %>% print(n = 12)
```

```
## # A tibble: 10 x 4
```

##		id	word	n	story
##		<int>	<chr>	<int>	<chr>
##	1	1	photograph	21	a_scandal_in_bohemia
##	2	1	door	13	a_scandal_in_bohemia
##	3	1	myself	11	a_scandal_in_bohemia
##	4	1	cried	10	a_scandal_in_bohemia
##	5	1	eyes	9	a_scandal_in_bohemia
##	6	1	lady	9	a_scandal_in_bohemia
##	7	1	heard	8	a_scandal_in_bohemia
##	8	1	indeed	8	a_scandal_in_bohemia
##	9	1	looked	8	a_scandal_in_bohemia
##	10	1	remarked	8	a_scandal_in_bohemia

Annotation engines



## Back ends

We have been able to get some real, interesting results by splitting our raw text into tokens. Some clever filtering and use of external datasets has gotten us some rough results in terms of character identification and the detection of themes.

To go deeper though, we need a more advanced natural language processing engine. These extract more granular features of the text, such as identifying parts of speech and tagging particular known entities.

## Back end, cont.

In **cleanNLP**, we currently provide back ends to two of the most well-known such libraries:

- ▶ **spaCy** a Python library primarily built for speed and stability
- ▶ **CoreNLP** a Java library built to have bleeding-edge functionality

## Initialising back ends

To use one of these backends in **cleanNLP**, simply run an alternative `init_` function before annotating the text. Either use:

```
library(cleanNLP)
init_spacy(model_name = "en")
anno <- run_annotators(paths)
nlp <- get_combine(anno)
```

Or:

```
library(cleanNLP)
init_coreNLP(language = "en")
anno <- run_annotators(paths)
nlp <- get_combine(anno)
```

## Annotation results

The resulting data set `nlp` also has one row per token, but now there are many additional features that have been learned from the text:

```
## # A tibble: 551,463 x 15
##       id    sid    tid word  lemma upos  pos      cid source relation
##   <int> <int> <int> <chr> <chr> <chr> <chr> <int> <int> <chr>
## 1     1     1     2 To    to    ADP   IN      1     0 ROOT
## 2     1     1     3 Sherl~ sherl~ PROPN NNP      4     4 compound
## 3     1     1     4 Holmes holmes PROPN NNP     13     2 pobj
## 4     1     1     5 she    -PRON- PRON  PRP     20     6 nsubj
## 5     1     1     6 is      be     VERB  VBZ     24     2 ccomp
## 6     1     1     7 always always ADV   RB     27     6 advmod
## # ... with 5.515e+05 more rows, and 5 more variables: word_source
## #   <chr>, lemma_source <chr>, entity_type <chr>, entity <chr>,
## #   spaces <int>
```

# Annotation tasks

NLP backends use models to learn features about the words and sentences in our raw text. Common tasks include:

- ▶ tokenisation
- ▶ lemmatisation
- ▶ sentence boundaries
- ▶ part of speech tags
- ▶ dependencies
- ▶ named entities
- ▶ coreferences
- ▶ sentiment analysis
- ▶ word embeddings

A collection of these running together (as they typically need to), is known as an **NLP Pipeline**. We will explain the meaning behind and some applications of many of these annotation tasks in these slides.

## Back end details

- ▶ options passed to the `init_` functions control which models and annotations are selected
- ▶ models have to be trained specifically for every natural language that they support
- ▶ more complex annotation tasks need to be trained separately for different styles of speech (i.e., Twitter versus Newspapers)
- ▶ libraries needed for these types of annotations require large external dependencies in order to run correctly
- ▶ in the interest of time, today we will simply provide the annotation objects for our corpora of study.

More detailed instructions for setting up either back end can be found on the `cleanNLP` repository and we are happy to help as best we can during the break or after the tutorial.

## Reading data

As mentioned above, we have already run the spaCy annotators on the corpus of Sherlock Holmes stories and made them available in the GitHub repository:

```
paths <- dir("data/holmes_stories", full.names = TRUE)
sh_meta <- data_frame(id = seq_along(paths),
                      story = stri_sub(basename(paths), 4, -5))
sh_nlp <- read_csv("data/sh_nlp.csv.gz")
```

## More accurate average sentence length


By using the sentence boundaries learned by the NLP pipeline, we can more accurately count the average length of the sentences in each text.

```
sh_nlp %>%  
  group_by(id, sid) %>%  
  mutate(sentence_end = tid == max(tid)) %>%  
  group_by(id) %>%  
  summarize(mean_sentence_len = n() / sum(sentence_end)) %>%  
  ggplot(aes(id, mean_sentence_len)) +  
    geom_line() +  
    geom_point()
```

Errors that might occur in our original method primarily include abbreviations such as “Dr.” and “S.O.S.”.



## More accurate average sentence length



# Lemmas

The most simple new column is the one titled `lemma`. This contains a reduced form of the token, for example converting all verbs into the same tense and all nouns into the singular case.

## Lemmas, examples

```
sh_nlp %>% filter(tolower(word) != lemma) %>%  
  select(word, lemma) %>% print(n = 10)
```

```
## # A tibble: 134,920 x 2  
##   word      lemma  
##   <chr>    <chr>  
## 1 she      -PRON-  
## 2 is       be  
## 3 I        -PRON-  
## 4 heard    hear  
## 5 him      -PRON-  
## 6 her      -PRON-  
## 7 his      -PRON-  
## 8 eyes     eye  
## 9 she      -PRON-  
## 10 eclipses eclipse  
## # ... with 1.349e+05 more rows
```

## Using lemmas

While minor, this assists with the topic discovery we did in the previous session by using the lemma frequency rather than the word frequency.

```
sh_lemmafr <- sh_nlp %>%  
  left_join(word_frequency, by = c("lemma" = "word")) %>%  
  filter(!is.na(frequency)) %>%  
  filter(frequency < 0.01) %>%  
  filter((tolower(word) == word)) %>%  
  count(id, lemma) %>%  
  group_by(id) %>%  
  top_n(n = 10, n) %>%  
  left_join(sh_meta, by = "id") %>%  
  arrange(id, desc(n))
```

## Using lemmas

```
sh_lemmafr %>% filter(id == 1) %>% print(n = 12)
```

```
## # A tibble: 10 x 4
##       id lemma          n story
##   <int> <chr>        <int> <chr>
## 1     1 photograph    21 a_scandal_in_bohemia
## 2     1 cry         15 a_scandal_in_bohemia
## 3     1 door        13 a_scandal_in_bohemia
## 4     1 minute      13 a_scandal_in_bohemia
## 5     1 eye         11 a_scandal_in_bohemia
## 6     1 hear        11 a_scandal_in_bohemia
## 7     1 lady        10 a_scandal_in_bohemia
## 8     1 rush        10 a_scandal_in_bohemia
## 9     1 gentleman    9 a_scandal_in_bohemia
## 10    1 throw        9 a_scandal_in_bohemia
```

## Using POS tags

Many of the tricks we used in the last session revolved around finding ways to approximate part of speech tags:

- ▶ stop words list, for example, removes (amongst other things) punctuation marks, pronouns, conjunctions, and interjections
- ▶ checking for upper case marks is really a hunt to identify proper nouns
- ▶ the frequency table is largely trying to remove verbs (there are far fewer of these and they tend to be more common), as well as common nouns

Proper part of speech tags can let us do these things more accurately as well as make other types of analysis possible.

## POS granularity

In primary or secondary school, you probably learned about a dozen or so parts of speech. These include nouns, verbs, adjectives, and so forth. Linguists in fact identify a far more granular set of part of speech tags, and even amongst themselves do not agree on a fixed set of such tags.

A commonly used one, and the one implemented by spaCy, are the Penn Treebank codes. These are given in our dataset under the pos variable.

# Penn Treebank

**Table 2**  
The Penn Treebank POS tagset.

1. CC	Coordinating conjunction	25. TO	<i>to</i>
2. CD	Cardinal number	26. UH	Interjection
3. DT	Determiner	27. VB	Verb, base form
4. EX	Existential <i>there</i>	28. VBD	Verb, past tense
5. FW	Foreign word	29. VBG	Verb, gerund/present participle
6. IN	Preposition/subordinating conjunction	30. VBN	Verb, past participle
7. JJ	Adjective	31. VBP	Verb, non-3rd ps. sing. present
8. JJR	Adjective, comparative	32. VBZ	Verb, 3rd ps. sing. present
9. JJS	Adjective, superlative	33. WDT	<i>wh</i> -determiner
10. LS	List item marker	34. WP	<i>wh</i> -pronoun
11. MD	Modal	35. WP\$	Possessive <i>wh</i> -pronoun
12. NN	Noun, singular or mass	36. WRB	<i>wh</i> -adverb
13. NNS	Noun, plural	37. #	Pound sign
14. NNP	Proper noun, singular	38. \$	Dollar sign
15. NNPS	Proper noun, plural	39. .	Sentence-final punctuation
16. PDT	Predeterminer	40. ,	Comma
17. POS	Possessive ending	41. :	Colon, semi-colon
18. PRP	Personal pronoun	42. (	Left bracket character
19. PP\$	Possessive pronoun	43. )	Right bracket character
20. RB	Adverb	44. "	Straight double quote
21. RBR	Adverb, comparative	45. '	Left open single quote
22. RBS	Adverb, superlative	46. "	Left open double quote
23. RP	Particle	47. '	Right close single quote
24. SYM	Symbol (mathematical or scientific)	48. "	Right close double quote



## Universal part of speech

Work has also been done to map these granular codes to language-agnostic codes known as universal parts of speech. Coincidentally, these universal parts of speech mimic those commonly taught in schools:

- ▶ *VERB*: verbs (all tenses and modes)
- ▶ *NOUN*: nouns (common and proper)
- ▶ *PRON*: pronouns
- ▶ *ADJ*: adjectives
- ▶ *ADV*: adverbs
- ▶ *ADP*: adpositions (prepositions and postpositions)
- ▶ *CONJ*: conjunctions
- ▶ *DET*: determiners
- ▶ *NUM*: cardinal numbers
- ▶ *PRT*: particles or other function words
- ▶ *X*: other: foreign words, typos, abbreviations
- ▶ *.*: punctuation

These are contained in the variable `upos`, and for today will be the most useful for our analysis.

## Top characters, again

Here, for example, is the analysis of key characters with our trick replaced by filtering on the proper noun tag “PROPN”:

```
sh_topchar <- sh_nlp %>%  
  filter(upos == "PROPN") %>%  
  count(id, word) %>%  
  group_by(id) %>%  
  top_n(n = 10, n) %>%  
  left_join(sh_meta, by = "id") %>%  
  arrange(id, desc(n))
```

## Top characters, again

```
sh_topchar %>% filter(id == 1) %>% print(n = Inf)
```

```
## # A tibble: 12 x 4
##       id word          n story
##   <int> <chr>      <int> <chr>
## 1     1  Holmes      48 a_scandal_in_bohemia
## 2     2  Majesty     18 a_scandal_in_bohemia
## 3     3  Irene       14 a_scandal_in_bohemia
## 4     4  Adler       13 a_scandal_in_bohemia
## 5     5  Briony      11 a_scandal_in_bohemia
## 6     6  King        11 a_scandal_in_bohemia
## 7     7  Lodge       11 a_scandal_in_bohemia
## 8     8  Sherlock    11 a_scandal_in_bohemia
## 9     9  Mr.         9  a_scandal_in_bohemia
## 10    10  Bohemia     7  a_scandal_in_bohemia
## 11    11  Norton      7  a_scandal_in_bohemia
## 12    12  Street      7  a_scandal_in_bohemia
```

## Compound words (optional)

A major shortcoming in our tabulation of proper nouns is that many of the proper nouns, in fact most in this case, are actually compound words. The proper way to analyse this data would be to collapse the compound words into a single combined token. It is relatively easy to do this in a slow way with loops. A fast, vectorized method with **dplyr** verbs is show in the code chunk below:

```
sh_compound <- sh_nlp %>%  
  filter(upos == "PROPN") %>%  
  group_by(id, sid) %>%  
  mutate(d = tid - lag(tid) - 1) %>%  
  mutate(d = ifelse(is.na(d), 1, d)) %>%  
  ungroup() %>%  
  mutate(d = cumsum(d)) %>%  
  group_by(d) %>%  
  summarize(id = first(id), sid = first(sid),  
            tid = first(tid),  
            thing = stri_c(word, collapse = " ")) %>%  
  select(-d) %>%  
  inner_join(sh_nlp, by = c("id", "sid", "tid"))
```

## Compound words (optional)

```
sh_compound %>% select(id, thing) %>% print(n = 10)
```

```
## # A tibble: 12,793 x 2
##       id thing
##   <int> <chr>
## 1     1  Sherlock Holmes
## 2     1  Irene Adler
## 3     1  Irene Adler
## 4     1   Holmes
## 5     1   Holmes
## 6     1 Baker Street
## 7     1  Odessa
## 8     1  Trepoff
## 9     1  Atkinson
## 10    1 Trincomalee
## # ... with 1.278e+04 more rows
```

# Entities

The task of finding characters, places, and other references to proper objects is common enough that it has been wrapped up into a specific annotation task known as named entity recognition (NER). Here are the first few entities from the annotation of our stories

```
results <- sh_nlp %>%  
  select(id, entity, entity_type) %>%  
  filter(!is.na(entity))
```

## Entities, cont.

```
results %>% print(n = 10)
```

```
## # A tibble: 18,939 x 3
##       id entity          entity_type
##   <int> <chr>          <chr>
## 1     1 1 Sherlock Holmes PERSON
## 2     2 1 Irene Adler   PERSON
## 3     3 1 one          CARDINAL
## 4     4 1 Grit          FAC
## 5     5 1 one          CARDINAL
## 6     6 1 one          CARDINAL
## 7     7 1 Irene Adler   PERSON
## 8     8 1 Holmes        PERSON
## 9     9 1 first        ORDINAL
## 10    10 1 Holmes        PERSON
## # ... with 1.893e+04 more rows
```

## NER characters

One benefit of this is that NER distinguishes between people and places, making our tabulation even more accurate:

```
sh_nerchar <- sh_nlp %>%  
  select(id, entity, entity_type) %>%  
  filter(!is.na(entity)) %>%  
  filter(entity_type == "PERSON") %>%  
  count(id, entity) %>%  
  group_by(id) %>%  
  top_n(n = 10, n) %>%  
  left_join(sh_meta, by = "id") %>%  
  arrange(id, desc(n))
```

```
sh_nerchar <- ungroup(sh_nerchar)
```



## NER characters, cont.

```
sh_nerchar %>% filter(id == 1) %>% print(n = Inf)
```

```
## # A tibble: 10 x 4
```

##		id	entity	n	story
##		<int>	<chr>	<int>	<chr>
##	1	1	Holmes	36	a_scandal_in_bohemia
##	2	1	Irene Adler	11	a_scandal_in_bohemia
##	3	1	Sherlock Holmes	9	a_scandal_in_bohemia
##	4	1	Briony Lodge	8	a_scandal_in_bohemia
##	5	1	Watson	6	a_scandal_in_bohemia
##	6	1	Godfrey Norton	4	a_scandal_in_bohemia
##	7	1	Your Majesty	4	a_scandal_in_bohemia
##	8	1	Temple	3	a_scandal_in_bohemia
##	9	1	Adler	2	a_scandal_in_bohemia
##	10	1	Count Von Kramm	2	a_scandal_in_bohemia

## Other entity categories

There are many other categories of named entities available within the spaCy and CoreNLP libraries, including:

- ▶ *ORGA*: Companies, agencies, institutions, etc.
- ▶ *MONEY*: Monetary values, including unit.
- ▶ *PERCENT*: Percentages.
- ▶ *DATE*: Absolute or relative dates or periods.
- ▶ *TIME*: Times smaller than a day.
- ▶ *NORP*: Nationalities or religious or political groups.
- ▶ *FACILITY*: Buildings, airports, highways, bridges, etc.
- ▶ *GPE*: Countries, cities, states.
- ▶ *LOC*: Non-GPE locations, mountain ranges, bodies of water.
- ▶ *PRODUCT*: Objects, vehicles, foods, etc. (Not services.)
- ▶ *EVENT*: Named hurricanes, battles, wars, sports events, etc.
- ▶ *WORK\_OF\_ART*: Titles of books, songs, etc.
- ▶ *LANGUAGE*: Any named language.
- ▶ *QUANTITY*: Measurements, as of weight or distance.
- ▶ *ORDINAL*: “first”, “second”, etc.
- ▶ *CARDINAL*: Numerals that do not fall under another type.

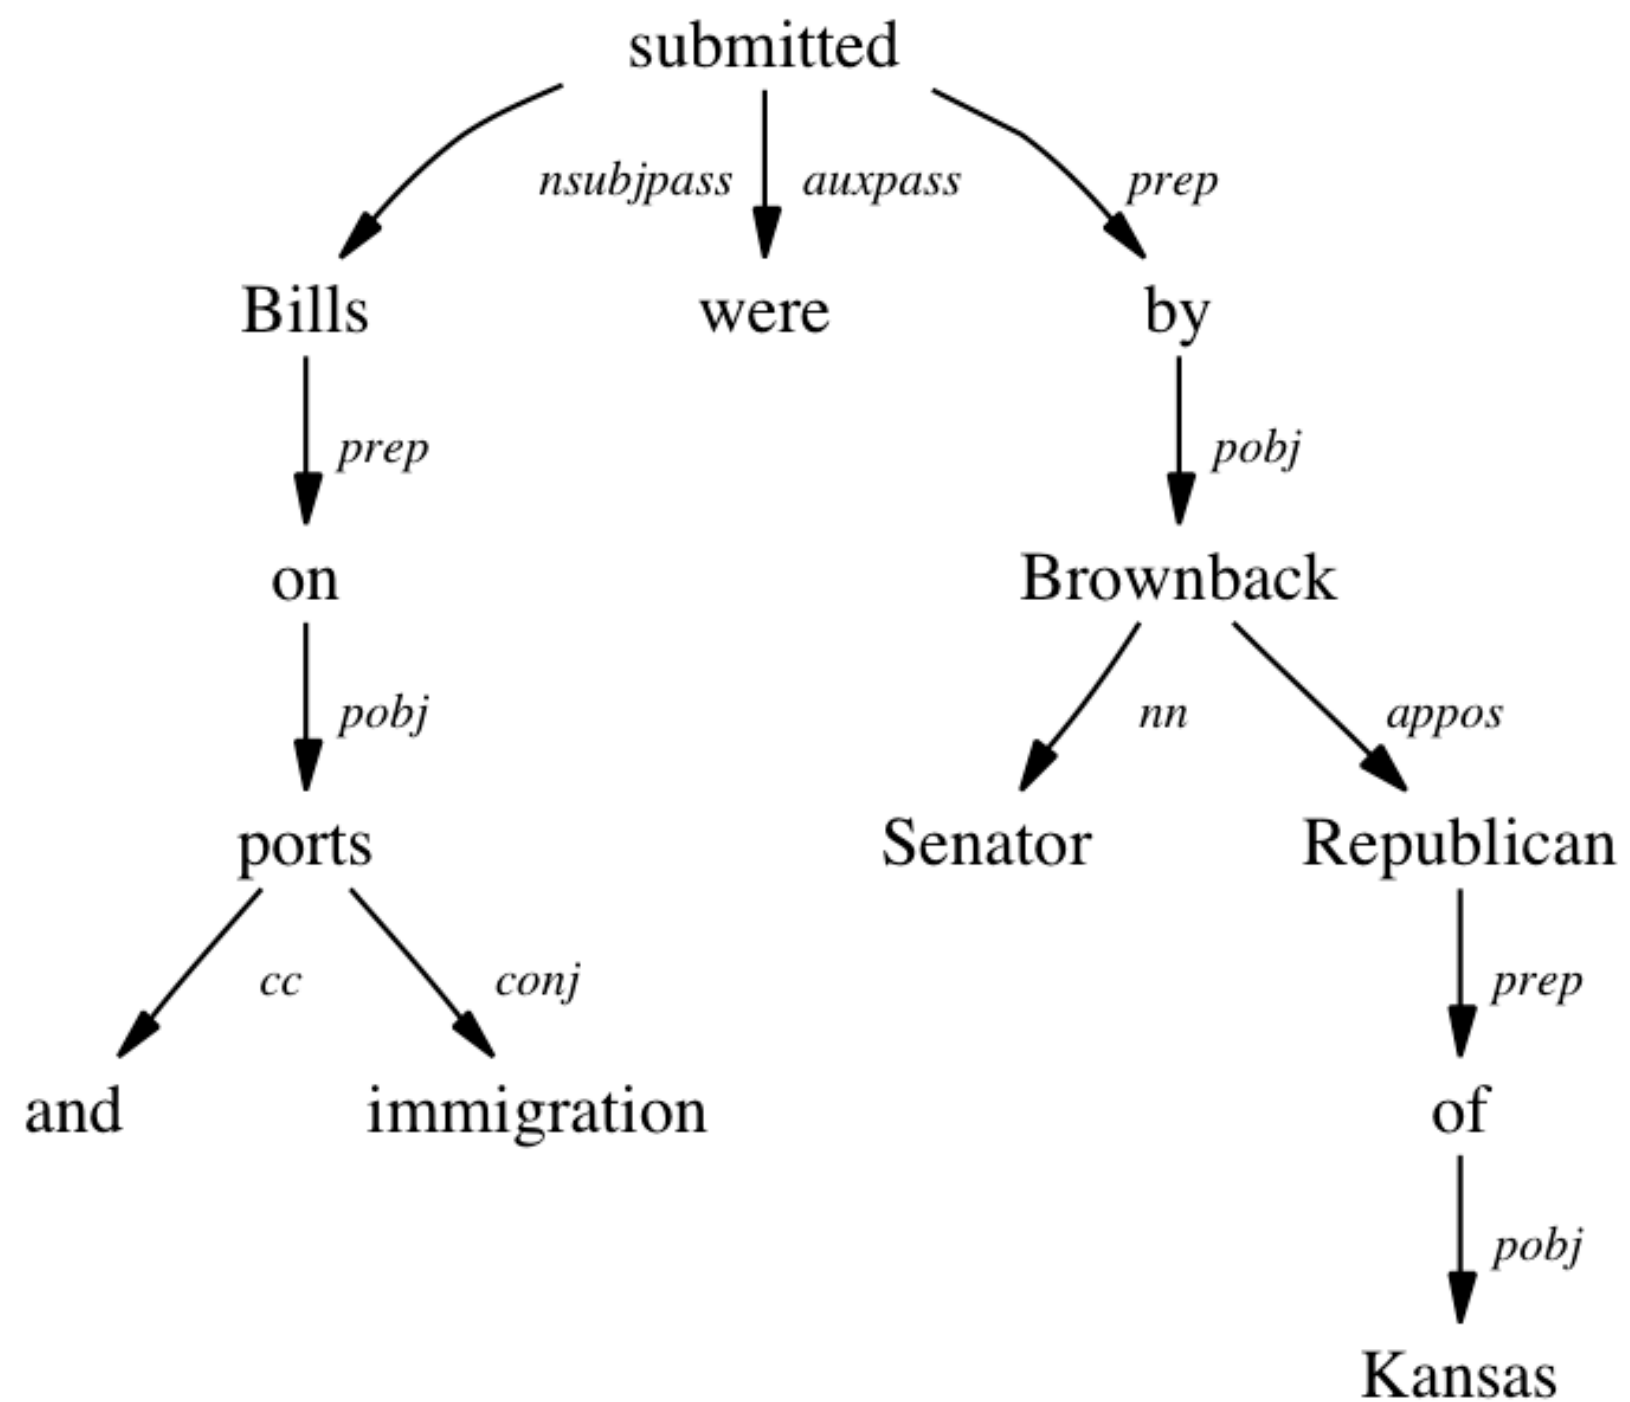
How might these be useful in various textual analyses?

# Dependencies

Dependencies are links between tokens within a sentence that indicate grammatical relationships.

For example, they link adjectives to the nouns they describe and adverbs to the verbs they modify. One of the most common dependencies is the direct object tag “dobj”, linking a verb to the noun that receives the action of the verb.

## Fully parsed sentence



## Dependencies, example

```
sh_nlp %>% filter(id == 1, sid == 1) %>%  
  select(word, source, relation, word_source)
```

```
## # A tibble: 9 x 4  
##   word      source relation word_source  
##   <chr>    <int> <chr>    <chr>  
## 1 To      0 ROOT     ROOT  
## 2 Sherlock 4 compound Holmes  
## 3 Holmes   2 pobj     To  
## 4 she      6 nsubj    is  
## 5 is       2 ccomp    To  
## 6 always   6 advmod   is  
## # ... with 3 more rows
```

## What are characters doing?

One way that dependencies can be useful is by determining which verbs are associated with each character by way of the 'nsubj' relation. Amongst other things, this can help identify sentiment, biases, and power dynamics.

In our corpus, we can use the 'nsubj' tag to identify verbs associated with our main characters:

```
sh_whatchar <- sh_nlp %>%  
  filter(relation == "nsubj") %>%  
  filter(upos == "PROPN") %>%  
  count(id, word, lemma_source) %>%  
  filter(n > 1)
```

## What are characters doing?

```
sh_whatchar %>% print(n = 12)
```

```
## # A tibble: 344 x 4
##       id word      lemma_source      n
##   <int> <chr>      <chr>      <int>
## 1     1 1 Holmes    murmur        2
## 2     2 1 Holmes    say           8
## 3     3 2 Holmes    remark        2
## 4     4 2 Holmes    say          11
## 5     5 2 I.        say           2
## 6     6 2 Merryweather be           2
## 7     7 2 Ross      be           3
## 8     8 2 Spaulding say           2
## 9     9 2 Wilson    be           2
## 10    10 2 Wilson    say           4
## 11    11 3 Angel     come          2
## 12    12 3 Holmes    remark        2
## # ... with 332 more rows
```

## Packages

To run the following code, we again make sure that the following packages are loaded (and installed) the following packages.

```
library(cleanNLP)
library(dplyr)
library(readr)
library(stringi)
library(ggplot2)
library(topicmodels)
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 3.4.4
```

```
library(ggrepel)
library(viridis)
```

```
## Warning: package 'viridis' was built under R version 3.4.4
```

```
library(magrittr)
theme_set(theme_minimal())
```

You will also need to download and set-up the tutorial's datasets.



# The Data

The President of the United States is constitutionally obligated to provide a report known as the 'State of the Union'. The report summarizes the current challenges facing the country and the president's upcoming legislative agenda.

We have run the spaCy NLP pipeline over this corpus and provide the output data in the GitHub repository.

```
sotu_nlp <- read_csv("data/sotu.csv.gz")  
sotu_meta <- read_csv("data/sotu_meta.csv")
```

## Sentence lengths

Just because we are doing text analysis is no excuse for not doing basic exploratory analysis of our data. What, for example, is the distribution of sentence lengths in the corpus?

```
sotu_nlp %>%  
  count(id, sid) %$%  
  quantile(n, seq(0,1,0.1))
```

##	0%	10%	20%	30%	40%	50%	60%	70%	80%	90%	100%
##	1	11	16	19	23	27	31	37	44	58	681

## Common nouns

What are the most common nouns in the corpus?

```
sotu_nlp %>%  
  filter(upos == "NOUN") %>%  
  count(lemma) %>%  
  top_n(n = 40, n) %>%  
  use_series(lemma)
```


##	[1]	"act"	"action"	"business"	"citizen"
##	[5]	"condition"	"country"	"duty"	"effort"
##	[9]	"force"	"government"	"interest"	"land"
##	[13]	"law"	"legislation"	"man"	"measure"
##	[17]	"nation"	"part"	"peace"	"people"
##	[21]	"policy"	"power"	"program"	"purpose"
##	[25]	"question"	"right"	"service"	"state"
##	[29]	"subject"	"system"	"tax"	"time"
##	[33]	"treaty"	"war"	"way"	"what"
##	[37]	"who"	"work"	"world"	"year"

## Length in words

Now, how long is each State of the Union in words? Does this differ based on whether it was given as a speech or a written document?

```
sotu_nlp %>%  
  count(id) %>%  
  group_by(id) %>%  
  left_join(sotu_meta, by = "id") %>%  
  ggplot(aes(year, n)) +  
    geom_line(color = grey(0.8)) +  
    geom_point(aes(color = sotu_type)) +  
    geom_smooth()
```

Length in words



## Summarising with dependencies

A straightforward way of extracting a high-level summary of the content of a speech is to extract all direct object object dependencies where the target noun is not a very common word.

Here is an example of this using the first address made by George W. Bush in 2001:

```
summary_2001 <- sotu_nlp %>%  
  left_join(sotu_meta, by = "id") %>%  
  filter(year == 2001, relation == "dobj") %>%  
  left_join(word_frequency, by = "word") %>%  
  filter(frequency < 0.001) %>%  
  select(id, word, word_source) %$%  
  sprintf("%s => %s", word_source, word)
```

# George W. Bush (2001)

summary\_2001

## [1] "take => oath"	"increasing => layoffs"
## [3] "buying => prescriptions"	"protects => trillion"
## [5] "makes => welcoming"	"accelerating => cleanup"
## [7] "fight => homelessness"	"allowing => taxpayers"
## [9] "provide => mentor"	"fight => illiteracy"
## [11] "promotes => compassion"	"end => profiling"
## [13] "stopping => abuses"	"pay => trillion"
## [15] "throw => darts"	"restores => fairness"
## [17] "restructure => defenses"	"promoting => internationalism"
## [19] "makes => downpayment"	"deploy => defenses"
## [21] "discard => relics"	"confronting => shortage"
## [23] "sound => footing"	"bridge => divides"
## [25] "minding => manners"	"divided => conscience"
## [27] "done => servants"	

## George W. Bush (2002)

```
head(summary_2002, 34)
```

```
## [1] "faces => dangers"      "urged => followers"
## [3] "brought => sorrow"    "found => diagrams"
## [5] "hold => hostages"     "eliminate => parasites"
## [7] "prevent => regimes"   "flaunt => hostility"
## [9] "develop => anthrax"   "kicked => inspectors"
## [11] "match => hatred"      "attack => allies"
## [13] "deploy => defenses"   "permit => regimes"
## [15] "increased => vigilance" "develop => vaccines"
## [17] "fight => anthrax"     "expand => patrols"
## [19] "track => arrivals"    "mean => neighborhoods"
## [21] "thank => attendants"  "defeat => recession"
## [23] "want => paycheck"     "set => posturing"
## [25] "reduce => dependency"  "offer => dignity"
## [27] "enact => safeguards"   "keeping => commitments"
## [29] "saw => selves"        "embracing => ethic"
## [31] "extending => compassion" "extend => compassion"
## [33] "await => knock"       "owns => aspirations"
```



## Woodrow Wilson (1919)

```
head(summary_1919, 34)
```

```
## [1] "save => inconvenience"      "produce => stagnation"
## [3] "produce => stagnation"      "made => interruption"
## [5] "keep => armies"           "-and => necessities"
## [7] "arrive => permitting"      "urge => necessity"
## [9] "produced => bitterness"     "remove => grievances"
## [11] "produces => dissatisfaction" "stir => disturbances"
## [13] "shown => willingness"      "bring => democratization"
## [15] "analyze => particulars"    "bid => pause"
## [17] "saps => vitality"          "treat => manifestations"
## [19] "touch => tissues"          "come => unrest"
## [21] "settle => disputes"         "devise => tribunal"
## [23] "lose => composure"         "realize => fruition"
```

## NLP and matrices

So far, we have done all of our analysis using a data frame where each token is given its own row. For modelling purposes, we often want to calculate the term frequency matrix. This matrix has one row per document and one column per unique token in the data set (although we can limit which tokens actually have a column).

Conveniently, **cleanNLP** provides the function `get_tfidf` for calculated this matrix.

## Document term frequency matrix

$$\begin{array}{l} \text{Text \#0001} \\ \text{Text \#0002} \\ \text{Text \#0003} \\ \vdots \\ \text{Text \#9500} \end{array} \begin{pmatrix} \text{I} & \text{and} & \dots & \text{commute} & \dots & \text{lol} \\ 20 & 55 & \dots & 0 & \dots & 0 \\ 34 & 72 & \dots & 5 & \dots & 0 \\ 6 & 34 & \dots & 0 & \dots & 4 \\ \vdots & \vdots & \dots & \vdots & \ddots & \vdots \\ 150 & 87 & \dots & 0 & \dots & 30 \end{pmatrix}$$

## get\_tfidf()

Here we will construct a term frequency matrix from only non-proper nouns:

```
sotu_tfidf <- sotu_nlp %>%  
  filter(pos %in% c("NN", "NNS")) %>%  
  get_tfidf(min_df = 0.05, max_df = 0.95,  
            type = "tfidf", tf_weight = "dnorm")
```

## NOTE: get\_tfidf has been renamed cnlp\_get\_tfidf

## NOTE: returning legacy output format from get\_tfidf

## get\_tfidf()

The output is a list with three elements: the term frequency inverse document frequency matrix, the ids of the documents corresponding to row names, and the vocabulary corresponding to the column names.

```
head(sotu_tfidf$vocab, 20)
```

```
## [1] "world"      "citizen"    "service"    "duty"
## [5] "system"     "right"      "man"        "program"
## [9] "policy"     "work"       "act"        "condition"
## [13] "subject"    "legislation" "force"      "effort"
## [17] "treaty"     "purpose"    "land"       "business"
```

```
head(sotu_tfidf$id, 20)
```

```
## [1] "1" "2" "3" "4" "5" "6" "7" "8" "9" "10" "11" "12"
## [13] "13" "14" "15" "16" "17" "18" "19" "20"
```

```
dim(sotu_tfidf$tfidf)
```

```
## [1] 236 2356
```

## PCA

What specifically can we do with this data? As a starting point, we will compute the principal components of the matrix. While base-R has great functions for doing this, we'll make use of the **cleanNLP** function `tidy_pca` which returns a data frame that makes plotting in **ggplot2** easier:

```
sotu_pca <- tidy_pca(sotu_tfidf$tfidf, sotu_meta)
```

```
## NOTE: tidy_pca has been renamed cnlp_pca
```

```
select(sotu_pca, president, party, PC1, PC2)
```


```
## # A tibble: 236 x 4
##   president      party      PC1    PC2
##   <chr>          <chr>    <dbl> <dbl>
## 1 George Washington Nonpartisan - 1.99  13.0
## 2 George Washington Nonpartisan - 4.83  16.7
## 3 George Washington Nonpartisan - 5.74  13.0
## 4 George Washington Nonpartisan - 3.34  12.2
## 5 George Washington Nonpartisan -16.9   18.7
## 6 George Washington Nonpartisan - 5.66  13.3
## # ... with 230 more rows
```

## PCA plot

While a simple scatter plot of this is easy to construct, we can tweak some of the default settings to get a really nice visualization of where each President's speeches cluster:

```
ggplot(sotu_pca, aes(PC1, PC2)) +  
  geom_point(aes(color = cut(year, 10, dig.lab = 4))) +  
  geom_text(data = filter(sotu_pca, !duplicated(president)))
```

PCA plot, cont.





# Topic Models

Topic models are a collection of statistical models for describing abstract themes within a textual corpus. Each theme is characterized by a collection of words that commonly co-occur; for example, the words 'crop', 'dairy', 'tractor', and 'hectare', might define a *farming* theme.

One of the most popular topic models is latent Dirichlet allocation (LDA), a Bayesian model where each topic is described by a probability distribution over a vocabulary of words. Each document is then characterized by a probability distribution over the available topics.

# LDA

To fit LDA on a corpus of text parsed by the **cleanNLP** package, the output of `get_tfidf` can be piped directly to the LDA function in the package **topicmodels**. The topic model function requires raw counts, so the type variable in `get_tfidf` is set to "tf".

```
sotu_tf <- sotu_nlp %>%  
  filter(pos %in% c("NN", "NNS")) %>%  
  get_tfidf(min_df = 0.05, max_df = 0.95,  
            type = "tf", tf_weight = "raw")  
tm <- LDA(sotu_tf$tf, k = 16, control = list(verbose = 1))
```

## NOTE: `get_tfidf` has been renamed `cnlp_get_tfidf`

## NOTE: returning legacy output format from `get_tfidf`

## Describing topics

We can describe each topic by giving the five most important words in each topic:

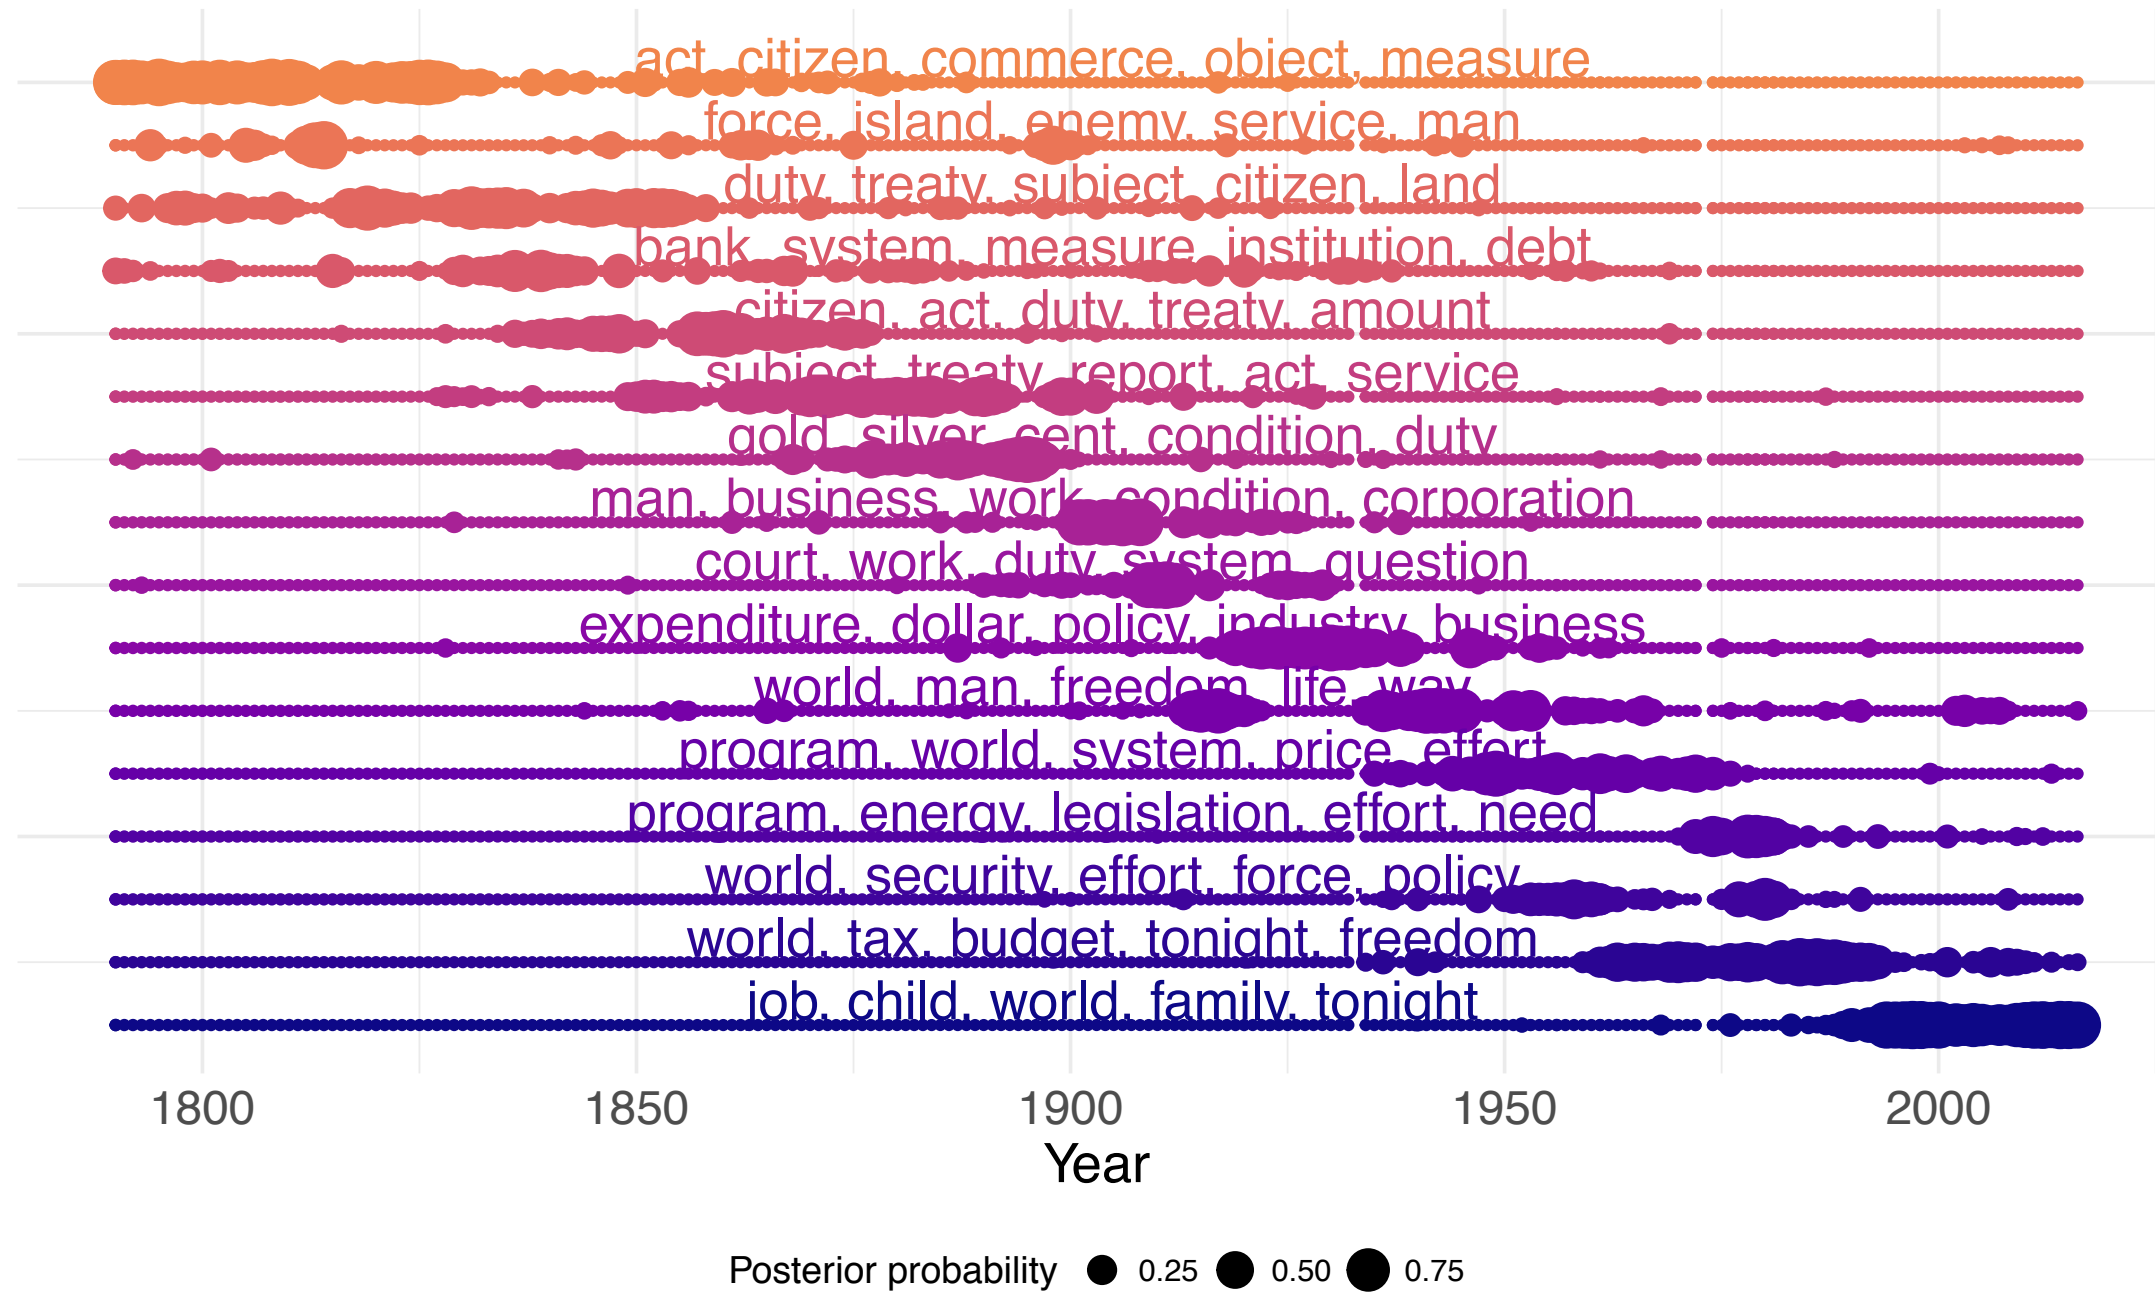
```
terms <- posterior(tm)$terms
topics <- posterior(tm)$topics
topic_df <- data_frame(topic = as.integer(col(topics)),
                      id = sotu_meta$id[as.integer(row(topics))],
                      val = as.numeric(topics)) %>%
  left_join(sotu_meta, by = "id")
top_terms <- apply(terms, 1,
  function(v) paste(sotu_tf$vocab[order(v,
    decreasing = TRUE)[1:5]], collapse = ", "))
top_terms <- as.character(top_terms)
```

## Describing topics

top\_terms

```
## [1] "act, citizen, commerce, object, measure"
## [2] "man, business, work, condition, corporation"
## [3] "duty, treaty, subject, citizen, land"
## [4] "world, man, freedom, life, way"
## [5] "bank, system, measure, institution, debt"
## [6] "program, world, system, price, effort"
## [7] "world, tax, budget, tonight, freedom"
## [8] "job, child, world, family, tonight"
## [9] "world, security, effort, force, policy"
## [10] "program, energy, legislation, effort, need"
## [11] "citizen, act, duty, treaty, amount"
## [12] "expenditure, dollar, policy, industry, business"
## [13] "court, work, duty, system, question"
## [14] "gold, silver, cent, condition, duty"
## [15] "force, island, enemy, service, man"
## [16] "subject, treaty, report, act, service"
```

# Topics over time



## Predictive models

A classifier that distinguishes speeches made by two presidents will be constructed here for the purpose of illustrating the topical and stylistic differences between them and their speech writers.

As a first step, a term-frequency matrix is extracted using the same technique as was used with the topic modeling function. However, here the frequency is computed for each sentence in the corpus rather than the document as a whole.

‘George Bush (2001-2008)’





## ‘Barack Obama (2009-2016)’





## Design matrix

The ability to do this seamlessly with a single additional mutate function defining a new id illustrates the flexibility of the get\_tfidf function.

```
df <- sotu_nlp %>%  
  left_join(sotu_meta, by = "id") %>%  
  filter(president %in% c("Barack Obama", "George W. Bush")) %>%  
  mutate(new_id = paste(id, sid, sep = "-")) %>%  
  filter(pos %in% c("NN", "NNS"))  
mat <- get_tfidf(df, min_df = 0, max_df = 1, type = "tf",  
                 tf_weight = "raw", doc_var = "new_id")
```

## NOTE: get\_tfidf has been renamed cnlp\_get\_tfidf

## NOTE: returning legacy output format from get\_tfidf

## Training and testing sets

It will be necessary to define a response variable  $y$  indicating whether this is a speech made by President Obama as well as a training flag indicating which speeches were made in odd numbered years.

```
m2 <- data_frame(new_id = mat$id) %>%  
  left_join(df[!duplicated(df$new_id),]) %>%  
  mutate(y = as.numeric(president == "Barack Obama")) %>%  
  mutate(train = (year %% 2 == 0))
```

## Elastic net

The output may now be used as input to the elastic net function provided by the **glmnet** package. This function fits a model of the form:

$$\beta = \operatorname{argmin}_b \left\{ ||y - Xb||_2 + \lambda \cdot (\alpha) ||b||_1 + \lambda \cdot (1 - \alpha) ||b||_2^2 \right\}$$

Cross-validation is used in order to select the best value of the model's tuning parameter  $\lambda$ .

```
model <- cv.glmnet(mat$tf[m2$train,], m2$y[m2$train],  
                  family = "binomial", alpha = 0.9)
```

The response is set to the binomial family given the binary nature of the response and training is trained on only those speeches occurring in odd-numbered years.

## Predicted probabilities

We can add the predicted probabilities to the dataset m2 with the following:

```
m2$pred <- predict(model, newx = mat$tf, type = "response",  
                  s = model$lambda.1se)  
select(m2, new_id, id, sid, president, year, pred)
```


```
## # A tibble: 4,821 x 6  
##   new_id    id  sid president      year  pred  
##   <chr>  <int> <int> <chr>      <int> <dbl>  
## 1 221-1    221     1 George W. Bush  2001 0.598  
## 2 221-2    221     2 George W. Bush  2001 0.636  
## 3 221-3    221     3 George W. Bush  2001 0.636  
## 4 221-4    221     4 George W. Bush  2001 0.784  
## 5 221-5    221     5 George W. Bush  2001 0.636  
## 6 221-6    221     6 George W. Bush  2001 0.784  
## # ... with 4,815 more rows
```

## Predicted probabilities

A boxplot of the predicted classes for each sentence within a speech is a good way of evaluating the model:

```
ggplot(m2, aes(factor(year), pred)) +  
  geom_boxplot(aes(fill = president))
```

# Predicted probabilities



## Model coefficients

One benefit of the penalized linear regression model is that it is possible to interpret the coefficients in a meaningful way. Here are the non-zero elements of the regression vector, coded as whether they have a positive (more Obama) or negative (more Bush) sign:

```
beta <- coef(model, s = model[["lambda"]][10])[-1]
sprintf("%s (%d)", mat$vocab, sign(beta))[beta != 0]
```

## [1]	"job (1)"	"nation (-1)"	"business (1)"
## [4]	"child (-1)"	"terrorist (-1)"	"freedom (-1)"
## [7]	"college (1)"	"company (1)"	"thing (1)"
## [10]	"peace (-1)"	"change (1)"	"enemy (-1)"
## [13]	"terror (-1)"	"hope (-1)"	"drug (-1)"
## [16]	"kid (1)"	"regime (-1)"	"class (1)"
## [19]	"industry (1)"	"member (-1)"	"relief (-1)"
## [22]	"liberty (-1)"	"compassion (-1)"	"enforcement (-1)"
## [25]	"medicine (-1)"	"death (-1)"	"11th (-1)"
## [28]	"homeland (-1)"	"will (-1)"	"character (-1)"
## [31]	"culture (-1)"		




‘Bill Clinton (1993-2000)’





# Predicted probabilities (Bush vs. Clinton)




## Model coefficients (Bush vs. Clinton)


```
beta <- coef(model, s = model[["lambda"]][9])[-1]
sprintf("%s (%d)", mat$vocab, sign(beta))[beta != 0]
```

## [1]	"year (1)"	"child (1)"	"family (1)"
## [4]	"care (1)"	"community (1)"	"freedom (-1)"
## [7]	"century (1)"	"parent (1)"	"thing (1)"
## [10]	"welfare (1)"	"terrorist (-1)"	"challenge (1)"
## [13]	"woman (-1)"	"crime (1)"	"terror (-1)"
## [16]	"enemy (-1)"	"hope (-1)"	"something (1)"
## [19]	"idea (1)"	"troop (-1)"	"gun (1)"
## [22]	"regime (-1)"	"relief (-1)"	"liberty (-1)"
## [25]	"danger (-1)"	"attack (-1)"	"institution (-1)"
## [28]	"compassion (-1)"	"coalition (-1)"	"dignity (-1)"
## [31]	"11th (-1)"	"oil (-1)"	"homeland (-1)"

‘George Washington (1789-1797)’



# Predicted probabilities (Washington vs. Clinton)



## Model coefficients (Washington vs. Clinton)

```
## [1] "provision (-1)"      "measure (-1)"      "object (-1)"
## [4] "attention (-1)"      "mean (-1)"         "session (-1)"
## [7] "consideration (-1)"  "militia (-1)"      "establishment (-1)"
## [10] "circumstance (-1)"   "commerce (-1)"     "satisfaction (-1)"
## [13] "tribe (-1)"          "commissioner (-1)" "execution (-1)"
## [16] "case (-1)"
```